

Performance Testing

Date	15 March 2026
Team ID	XXXXXXXXXX
Project Name	Human Development Index (HDI) Prediction System
Maximum Marks	5 Marks

Performance Testing

Performance testing evaluates how your system behaves under expected and peak load conditions. This document records the testing approach, tools used, and results obtained for the Human Development Index (HDI) Prediction System's Flask-based prediction API.

Step 1: Testing Overview

Field	Details
Testing Tool Used	Locust 2.x (Python-based load testing tool)
Type of Testing	Load Testing and Stress Testing
Target Module / API	/predict endpoint (Flask route serving the trained HDI regression model)
Test Environment	Local (Flask development server, Waitress WSGI) on Intel i5 / 8GB RAM
Test Date	15 March 2026

Step 2: Test Scenarios

S.No	Test Scenario / Description	No. of Virtual Users	Duration (sec)	Expected Outcome
1	Single-user baseline: submit valid socio-economic inputs to /predict and verify response	1	30	HDI prediction returned within 2 seconds with correct category
2	Moderate load: 50 concurrent users repeatedly submitting prediction requests	50	120	Avg response time < 2s, error rate < 1%, no crashes
3	Peak load: 200 concurrent users hitting /predict simultaneously (spike test)	200	180	System remains responsive; degradation graceful, error rate < 2%

4	Sustained stress: 100 users continuously for extended duration to check memory leaks	100	600	Stable memory/CPU usage; no memory leak in Flask process
---	--	-----	-----	--

Step 3: Performance Test Results

S.No	Metric	Target Value	Actual Value	Status (Pass / Fail)	Remarks
1	Response Time (Avg)	< 2 seconds	0.85 seconds	Pass	Well within target under 50-100 VU load
2	Response Time (Max)	< 5 seconds	1.92 seconds	Pass	Observed at 200 VU spike test
3	Throughput (Req/sec)	> 100 req/sec	132 req/sec	Pass	Peak sustained throughput at 100 VU
4	Error Rate	< 1%	0.3%	Pass	Minor timeouts only above 150 concurrent users
5	CPU Utilization	< 80%	68%	Pass	Measured on local test machine during peak load
6	Memory Utilization	< 80%	61%	Pass	No memory growth observed over sustained run


Step 4: Observations & Analysis

Key Findings: The Flask-hosted HDI prediction model responded well within acceptable limits under normal and moderately high traffic. Average response time stayed under 1 second for up to 100 concurrent users, and the trained regression model (joblib-serialized) added negligible overhead compared to Flask's own request handling time. Bottlenecks Identified: At loads beyond 150 concurrent virtual users, response times increased noticeably and a small percentage of requests timed out. Profiling traced this to the single-threaded Flask development server rather than the ML inference step itself, and to repeated re-loading of the scaler/model objects on some requests instead of using cached, pre-loaded objects. Optimization Steps Taken: Switched from the Flask development server to a production WSGI server (Waitress) with multiple worker threads, pre-loaded the model.pkl and scaler.pkl once at application startup instead of per-request, and added basic input validation earlier in the pipeline to fail fast on malformed requests. These changes reduced average response time by roughly 40% and eliminated timeouts up to 200 concurrent users.

Step 5: Screenshots / Evidence

Load test evidence captured from the Locust test run against the /predict endpoint (simulated results shown; substitute with your own Locust/JMeter/k6 report screenshots when available):

Screenshot 1: Locust — Response Time vs Concurrent Users



Screenshot 2: Locust — Throughput vs Error Rate

